Name: Manasi. B. Kshirsagar

PRN: S17111009

Roll No.: 06

Class: BE Comp SS

Code (Maximum):

```
%%cu
#include <cstdio>
#include <iostream>
using namespace std;

__global__ void maxi(int* a, int* b, int n)
{
 int block = 256 * blockIdx.x;
 int max = 0;
 for (int i = block; i < min(256 + block, n); i++)
 {
 if (max < a[i])
 {
 max = a[i];
 }
 }
 b[blockIdx.x] = max;
}
```

```
int main()
{
 int n;
 n = 200;
 int a[n];
 for (int i = 0; i < n; i++)
 {
 a[i] = rand() % n;
 cout << a[i] << " ";
 }
 cout<<"\n";
 cudaEvent_t start, end;
 int *ad, *bd;
 int size = n * sizeof(int);
 cudaMalloc(&ad, size);
 cudaMemcpy(ad, a, size, cudaMemcpyHostToDevice);
 int grids = ceil(n * 1.0f / 256.0f);
 cudaMalloc(&bd, grids * sizeof(int));

 dim3 grid(grids, 1);
 dim3 block(1, 1);

 cudaEventCreate(&start);
 cudaEventCreate(&end);
 cudaEventRecord(start);
```

```
 while (n > 1)
 {
 maxi<<<grids, block>>>(ad, bd, n);
 n = ceil(n * 1.0f / 256.0f);
 cudaMemcpy(ad, bd, n * sizeof(int),
cudaMemcpyDeviceToDevice);
 }
 cudaEventRecord(end);
 cudaEventSynchronize(end);
 float time = 0;
 cudaEventElapsedTime(&time, start, end);
 int ans[2];
 cudaMemcpy(ans, ad, 4, cudaMemcpyDeviceToHost);
 cout<<"The maximum element is : " << ans[0] << endl;
 cout<<"The time required : ";
 cout<<time<<endl;
 return 0;
}
```

Output:

Code (Minimum):

```
%%cu
#include <cstdio>
#include <iostream>
using namespace std;

__global__ void mini(int* a, int* b, int n)
{
 int block = 256 * blockIdx.x;
 int minimum = 0;
 for (int i = block; i < min(256 + block, n); i++)
 {
 if (minimum > a[i])
 {
 minimum = a[i];
```

```
 }
 }
 b[blockIdx.x] = minimum;
}

int main()
{
 int n;
 n = 200;
 int a[n];
 for (int i = 0; i < n; i++)
 {
 a[i] = rand() % n;
 cout << a[i] << " ";
 }
 cout<<"\n";
 cudaEvent_t start, end;
 int *ad, *bd;
 int size = n * sizeof(int);
 cudaMalloc(&ad, size);
 cudaMemcpy(ad, a, size, cudaMemcpyHostToDevice);
 int grids = ceil(n * 1.0f / 256.0f);
 cudaMalloc(&bd, grids * sizeof(int));

 dim3 grid(grids, 1);
 dim3 block(1, 1);
```

```
 cudaEventCreate(&start);
 cudaEventCreate(&end);
 cudaEventRecord(start);

 while (n > 1)
 {
 mini<<<grids, block>>>(ad, bd, n);
 n = ceil(n * 1.0f / 256.0f);
 cudaMemcpy(ad, bd, n * sizeof(int),
cudaMemcpyDeviceToDevice);
 }
 cudaEventRecord(end);
 cudaEventSynchronize(end);
 float time = 0;
 cudaEventElapsedTime(&time, start, end);
 int ans[2];
 cudaMemcpy(ans, ad, 4, cudaMemcpyDeviceToHost);
 cout<<"The minimum element is : " << ans[0] << endl;
 cout<<"The time required : ";
 cout<<time<<endl;
 return 0;
}
```

Output:

Code (Std Dev & Variance):

```
%%cu
#include<iostream>
#include<cstdio>
using namespace std;
__global__ void var(int *a,int *b,int n,float mean)
{
 int block=256*blockIdx.x;
 float sum=0;
 for(int i=block;i<min(block+256,n);i++)
 {
 sum=sum+(a[i]-mean)*(a[i]-mean);
 }
 b[blockIdx.x]=sum;
```

```
}
__global__ void sum(int *a,int *b,int n)
{
 int block=256*blockIdx.x;
 int sum=0;
 for(int i=block;i<min(block+256,n);i++)
 {
 sum=sum+a[i];
 }
 b[blockIdx.x]=sum;
}
int main()
{
int n;
n=200;
int a[n];
cout<<"Elements: ";
for(int i=0;i<n;i++)
{
a[i]=rand()%n;
cout<<a[i]<<" ";
}
int *ad,*bd;
int size=n*sizeof(int);
cudaMalloc(&ad,size);
cudaMemcpy(ad,a,size,cudaMemcpyHostToDevice);
```

```
int grids=ceil(n*1.0f/256.0f);
cudaMalloc(&bd,grids*sizeof(int));
dim3 grid(grids,1);
dim3 block(1,1);
int p=n;
cudaEvent_t start,end;
cudaEventCreate(&start);
cudaEventCreate(&end);
cudaEventRecord(start);
while(n>1)
{
 sum<<<grid,block>>>(ad,bd,n);
 n=ceil(n*1.0f/256.0f);

cudaMemcpy(ad,bd,n*sizeof(int),cudaMemcpyDeviceToDevice);
}
cudaEventRecord(end);
cudaEventSynchronize(end);
float time=0;
cudaEventElapsedTime(&time,start,end);
cout<<endl<<"The Time is "<<time<<endl;
int add[2];
n=p;
cudaMemcpy(add,ad,4,cudaMemcpyDeviceToHost);
float mean=0.0f;
mean=add[0]/(n*1.0f);
```

```
cout<<"The Mean is "<<mean<<endl;
cudaMalloc(&ad,size);
cudaMemcpy(ad,a,size,cudaMemcpyHostToDevice);
cudaMalloc(&bd,grids*sizeof(int));
var<<<grid,block>>>(ad,bd,n,mean);
n=ceil(n*1.0f/256.0f);
sum<<<grid,block>>>(bd,ad,n);
cudaMemcpy(add,ad,4,cudaMemcpyDeviceToHost);
float sd=sqrt(add[0]/p*1.0f);
cout<<"The Standard Deviation is "<<sd<<endl;
}
```

Output: