

Name: Manasi. B. Kshirsagar

PRN: ST7111009

Roll No.: 06

Class: BE Comp SS

Code:

```
#include<iostream>
```

```
#include<cstdio>
```

```
#include<cstdlib>
```

```
#include<cuda_runtime.h>
```

```
using namespace std;
```

```
__global__ void minimum(int *input)
```

```
{
```

```
    int tid=threadIdx.x;
```

```
    auto step_size=1;
```

```
    int number_of_threads=blockDim.x;
```

```
    while(number_of_threads>0)
```

```
    {
```

```
        if(tid<number_of_threads)
```

```
        {
```

```
            int first=tid*step_size*2;
```

```
            int second=first+step_size;
```

```

        if(input[second]<input[first])
            input[first]=input[second];
    }
    step_size=step_size*2;
    number_of_threads/=2;
}
}

```

```

__global__ void max(int *input)
{
    int tid=threadIdx.x;
    auto step_size=1;
    int number_of_threads=blockDim.x;
    while(number_of_threads>0)
    {
        if(tid<number_of_threads)
        {
            int first=tid*step_size*2;
            int second=first+step_size;
            if(input[second]>input[first])
                input[first]=input[second];
        }
        step_size*=2;
        number_of_threads/=2;
    }
}

```

```

__global__ void sum(int *input)
{
    const int tid=threadIdx.x;
    auto step_size=1;
    int number_of_threads=blockDim.x;
    while(number_of_threads>0)
    {
        if(tid<number_of_threads)
        {
            const int first=tid*step_size*2;
            const int second=first+step_size;
            input[first]=input[first]+input[second];
        }
        step_size = step_size*2;;
        number_of_threads =number_of_threads/2;
    }
}

```

__global__ void average(int *input) //You can use above sum() to calculate sum and divide it by num_of_elements

```

{
    const int tid=threadIdx.x;
    auto step_size=1;
    int number_of_threads=blockDim.x;
    int totalElements=number_of_threads*2;
    while(number_of_threads>0)

```

```

{
    if(tid<number_of_threads)
    {
        const int first=tid*step_size*2;
        const int second=first+step_size;
        input[first]=input[first]+input[second];
    }
    step_size = step_size*2;;
    number_of_threads =number_of_threads/2;
}
input[0]=input[0]/totalElements;
}

```

```

int main()
{
    int n;
    n=200;
    srand(n);
    int *arr=new int[n];
    int min=20000; //Any Large Number would work
    cout<<"Elements are: "<<endl;
    //# Generate Input array using rand()
    for(int i=0;i<n;i++)
    {
        arr[i]=rand()%n;
        if(arr[i]<min)

```

```

        min=arr[i];
        cout<<arr[i]<<" ";
    }

    int size=n*sizeof(int); //calculate no. of bytes for array
    int *arr_d,result1;

    /// Allocate memory for min Operation
    cudaMalloc(&arr_d,size);
    cudaMemcpy(arr_d,arr,size,cudaMemcpyHostToDevice);
    minimum<<<1,n/2>>>(arr_d);
    cudaMemcpy(&result1,arr_d,sizeof(int),cudaMemcpyDeviceToHost);
    cout<<endl<<"The minimum element is "<<result1<<endl;
    cout<<"The min element (using CPU) is "<<min<<endl;

    ///MAX OPERATION
    int *arr_max,maxValue;
    cudaMalloc(&arr_max,size);
    cudaMemcpy(arr_max,arr,size,cudaMemcpyHostToDevice);
    max<<<1,n/2>>>(arr_max);
    cudaMemcpy(&maxValue,arr_max,sizeof(int),cudaMemcpyDeviceToHost);
    cout<<"The maximum element is "<<maxValue<<endl;

    ///SUM OPERATION
    int *arr_sum,sumValue;
    cudaMalloc(&arr_sum,size);

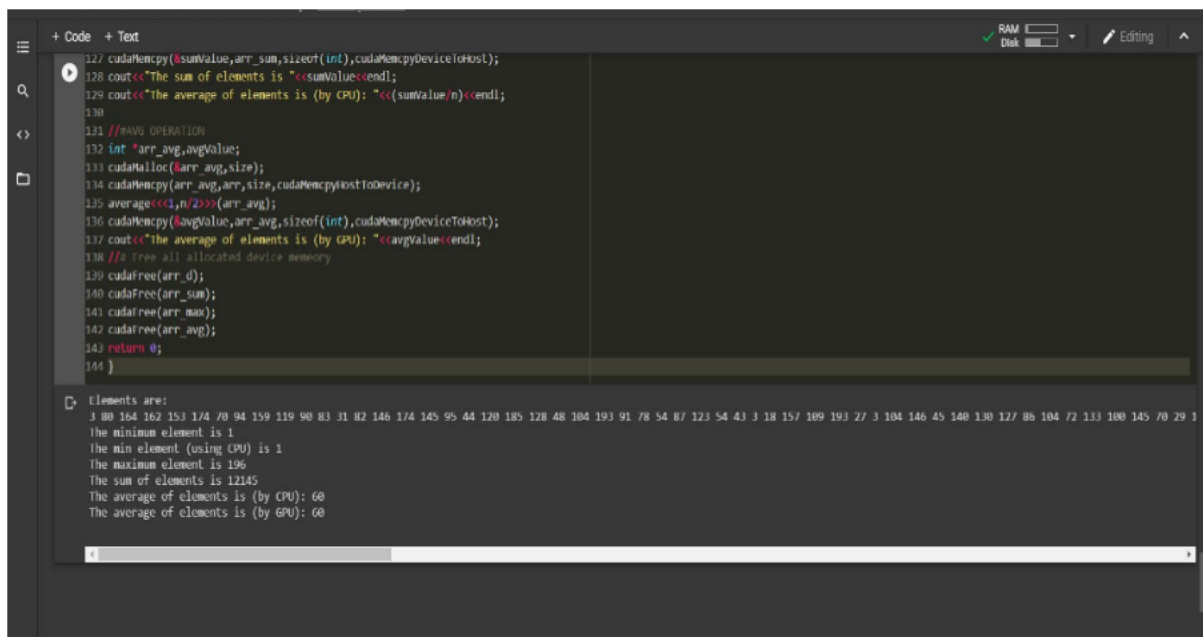
```

```
cudaMemcpy(arr_sum,arr,size,cudaMemcpyHostToDevice);
sum<<<1,n/2>>>(arr_sum);
cudaMemcpy(&sumValue,arr_sum,sizeof(int),cudaMemcpyDeviceToHost);
cout<<"The sum of elements is "<<sumValue<<endl;
cout<<"The average of elements is (by CPU): "<<(sumValue/n)<<endl;
```

```
///AVG OPERATION
```

```
int *arr_avg,avgValue;
cudaMalloc(&arr_avg,size);
cudaMemcpy(arr_avg,arr,size,cudaMemcpyHostToDevice);
average<<<1,n/2>>>(arr_avg);
cudaMemcpy(&avgValue,arr_avg,sizeof(int),cudaMemcpyDeviceToHost);
cout<<"The average of elements is (by GPU): "<<avgValue<<endl;
///Free all allocated device memory
cudaFree(arr_d);
cudaFree(arr_sum);
cudaFree(arr_max);
cudaFree(arr_avg);
return 0;
}
```

Output:



The screenshot shows a C++ IDE with a code editor and an output window. The code editor contains the following C++ code:

```
127 cudaMemcpy(&sumValue, arr_sum, sizeof(int), cudaMemcpyDeviceToHost);
128 cout<<"The sum of elements is "<<sumValue<<endl;
129 cout<<"The average of elements is (by CPU): "<<(sumValue/n)<<endl;
130
131 //AUG OPERATION
132 int *arr_avg, avgValue;
133 cudaMalloc(&arr_avg, size);
134 cudaMemcpy(arr_avg, arr, size, cudaMemcpyHostToDevice);
135 average<<1, n/2>>(arr_avg);
136 cudaMemcpy(&avgValue, arr_avg, sizeof(int), cudaMemcpyDeviceToHost);
137 cout<<"The average of elements is (by GPU): "<<avgValue<<endl;
138 //A free all allocated device memory
139 cudaFree(arr_d);
140 cudaFree(arr_sum);
141 cudaFree(arr_max);
142 cudaFree(arr_avg);
143 return 0;
144 }
```

The output window displays the following results:

```
Elements are:
3 80 164 162 153 174 70 94 159 119 90 83 31 82 146 174 145 95 44 128 185 128 48 104 193 91 78 54 87 123 54 43 3 18 157 169 193 27 3 104 146 45 140 130 127 86 104 72 133 100 145 70 29 1
The minimum element is 1
The min element (using CPU) is 1
The maximum element is 196
The sum of elements is 12145
The average of elements is (by CPU): 60
The average of elements is (by GPU): 60
```